Autodiff
Automatic Differentiation C++ Library

Matthew Pulver

March 9, 2019

1 Synopsis

```
#include <boost/math/differentiation/autodiff.hpp>

namespace boost { namespace math { namespace differentiation {

// Type for variables and constants.
template<typename RealType, size_t Order, size_t... Orders>
using autodiff_fvar = typename detail::nest_fvar<RealType,Order,Orders...>::type;

// Function returning a variable of differentiation.
template<typename RealType, size_t Order, size_t... Orders>
autodiff_fvar<RealType,Order,Orders...> make_fvar(const RealType& ca);

// Type of combined autodiff types.
template<typename RealType, typename... RealTypes>
using promote = typename detail::promote_args_n<RealType,RealTypes...>::type;

namespace detail {

// Single autodiff variable. Independent variables are created by nesting.
template<typename RealType, size_t Order>
class fvar
{
  public:

    // Query return value of function to get the derivatives.
    template<typename... Orders>
    get_type_at<RealType, sizeof...(Orders)-1> derivative(Orders... orders) const;

    // All of the arithmetic and comparison operators are overloaded.
    template<typename RealType2, size_t Order2>
    fvar& operator+=(const fvar<RealType2,Order2>&);

    fvar& operator+=(const root_type&);

    // ...
};

// Standard math functions are overloaded and called via argument-dependent lookup (ADL).
template<typename RealType, size_t Order>
fvar<RealType,Order> floor(const fvar<RealType,Order>&);

template<typename RealType, size_t Order>
fvar<RealType,Order> exp(const fvar<RealType,Order>&);

// ...

} // namespace detail

} } } // namespace boost::math::differentiation
```

2 Description

Autodiff is a header-only C++ library that facilitates the automatic differentiation (forward mode) of mathematical functions of single and multiple variables.

This implementation is based upon the Taylor series expansion of an analytic function f at the point x_0:

$$f(x_0 + \varepsilon) = f(x_0) + f'(x_0)\varepsilon + \frac{f''(x_0)}{2!}\varepsilon^2 + \frac{f'''(x_0)}{3!}\varepsilon^3 + \cdots$$

$$= \sum_{n=0}^{N} \frac{f^{(n)}(x_0)}{n!}\varepsilon^n + O\left(\varepsilon^{N+1}\right).$$

The essential idea of autodiff is the replacement of numbers with polynomials in the evaluation of f. By inputting the first-order polynomial $x_0 + \varepsilon$, the resulting polynomial in ε contains the function's derivatives within the coefficients. Each coefficient is equal to a derivative of its respective order, divided by the factorial of the order.

Assume one is interested in calculating the first N derivatives of f at x_0. Then without any loss of precision to the calculation of the derivatives, all terms $O\left(\varepsilon^{N+1}\right)$ that include powers of ε greater than N can be discarded, and under these truncation rules, f provides a polynomial-to-polynomial transformation:

$$f \quad : \quad x_0 + \varepsilon \quad \mapsto \quad \sum_{n=0}^{N} y_n \varepsilon^n = \sum_{n=0}^{N} \frac{f^{(n)}(x_0)}{n!}\varepsilon^n.$$

C++'s ability to overload operators and functions allows for the creation of a class `fvar` (forward-mode autodiff variable) that represents polynomials in ε. Thus the same algorithm that calculates the numeric value of $y_0 = f(x_0)$ is also used to calculate the polynomial $\sum_{n=0}^{N} y_n \varepsilon^n = f(x_0 + \varepsilon)$. The derivatives are then found from the product of the respective factorial and coefficient:

$$\frac{d^n f}{dx^n}(x_0) = n! y_n.$$

3 Examples

3.1 Example 1: Single-variable derivatives

3.1.1 Calculate derivatives of $f(x) = x^4$ at $x = 2$.

In this example, `autodiff_fvar<double,5>` is a data type that can hold a polynomial of up to degree 5, and the `make_fvar<double,5> x(2)` represents the polynomial $2+\varepsilon$. Internally, this is modeled by a `std::array<double,6>` whose elements {2, 1, 0, 0, 0, 0} correspond to the 6 coefficients of the polynomial upon initialization. Its fourth power is a polynomial with coefficients y = {16, 32, 24, 8, 1, 0}. The derivatives are obtained using the formula $f^{(n)}(2) = n! * \text{y[n]}$.

```
#include <boost/math/differentiation/autodiff.hpp>
#include <iostream>

template<typename T>
T fourth_power(const T& x)
{
    const T second_power = x * x;
    return second_power * second_power;
}

int main()
{
    using namespace boost::math::differentiation;

    constexpr int Order=5; // The highest order derivative to be calculated.
    const autodiff_fvar<double,Order> x = make_fvar<double,Order>(2); //Find derivatives at x=2.
```

```
    const autodiff_fvar<double,Order> y = fourth_power(x);
    for (int i=0 ; i<=Order ; ++i)
        std::cout << "y.derivative("<<i<<") = " << y.derivative(i) << std::endl;
    return 0;
}
/* Output:
y.derivative(0) = 16
y.derivative(1) = 32
y.derivative(2) = 48
y.derivative(3) = 48
y.derivative(4) = 24
y.derivative(5) = 0
*/
```

The above calculates

$$
\begin{aligned}
\texttt{y.derivative(0)} =& \quad f(2) = & x^4\big|_{x=2} &= 16 \\
\texttt{y.derivative(1)} =& \quad f'(2) = & 4 \cdot x^3\big|_{x=2} &= 32 \\
\texttt{y.derivative(2)} =& \quad f''(2) = & 4 \cdot 3 \cdot x^2\big|_{x=2} &= 48 \\
\texttt{y.derivative(3)} =& \quad f'''(2) = & 4 \cdot 3 \cdot 2 \cdot x|_{x=2} &= 48 \\
\texttt{y.derivative(4)} =& f^{(4)}(2) = & 4 \cdot 3 \cdot 2 \cdot 1 &= 24 \\
\texttt{y.derivative(5)} =& f^{(5)}(2) = & & 0
\end{aligned}
$$

3.2 Example 2: Multi-variable mixed partial derivatives with multi-precision data type

3.2.1 Calculate $\frac{\partial^{12} f}{\partial w^3 \partial x^2 \partial y^4 \partial z^3}(11, 12, 13, 14)$ with a precision of about 50 decimal digits, where $f(w, x, y, z) = \exp\left(w \sin\left(\frac{x \log(y)}{z}\right) + \sqrt{\frac{wz}{xy}}\right) + \frac{w^2}{\tan(z)}$.

In this example, the data type `autodiff_fvar<cpp_bin_float_50,Nw,Nx,Ny,Nz>` represents a multivariate polynomial in 4 independent variables, where the highest powers of each are `Nw`, `Nx`, `Ny` and `Nz`. The underlying arithmetic data type, referred to as `root_type`, is `boost::multiprecision::cpp_bin_float_50`. The internal data structure is `std::array<std::array<std::array<std::array<cpp_bin_float_50,Nz+1>,Ny+1>,Nx+1>,Nw+1>`. In general, the `root_type` is always the first template parameter to `autodiff_fvar<>` followed by the maximum derivative order that is to be calculated for each independent variable.

When variables are initialized with `make_var<...>()`, the position of the last derivative order given in the template parameter pack determines which variable is taken to be independent. In other words, it determines which of the 4 different polynomial variables $\varepsilon_w, \varepsilon_x, \varepsilon_y$, or ε_z are to be added to the constant term:

$$
\begin{aligned}
\texttt{make\_fvar<cpp\_bin\_float\_50,Nw>(11)} &= 11 + \varepsilon_w \\
\texttt{make\_fvar<cpp\_bin\_float\_50,0,Nx>(12)} &= 12 + \varepsilon_x \\
\texttt{make\_fvar<cpp\_bin\_float\_50,0,0,Ny>(13)} &= 13 + \varepsilon_y \\
\texttt{make\_fvar<cpp\_bin\_float\_50,0,0,0,Nz>(14)} &= 14 + \varepsilon_z
\end{aligned}
$$

Instances of different types are automatically promoted to the smallest multi-variable type that accommodates both when they are arithmetically combined (added, subtracted, multiplied, divided.)

```
#include <boost/math/differentiation/autodiff.hpp>
#include <boost/multiprecision/cpp_bin_float.hpp>
#include <iostream>

template<typename T>
T f(const T& w, const T& x, const T& y, const T& z)
```

3

```
{
  using namespace std;
  return exp(w*sin(x*log(y)/z) + sqrt(w*z/(x*y))) + w*w/tan(z);
}

int main()
{
  using cpp_bin_float_50 = boost::multiprecision::cpp_bin_float_50;
  using namespace boost::math::differentiation;

  constexpr int Nw=3; // Max order of derivative to calculate for w
  constexpr int Nx=2; // Max order of derivative to calculate for x
  constexpr int Ny=4; // Max order of derivative to calculate for y
  constexpr int Nz=3; // Max order of derivative to calculate for z
  using var = autodiff_fvar<cpp_bin_float_50,Nw,Nx,Ny,Nz>;
  const var w = make_fvar<cpp_bin_float_50,Nw>(11);
  const var x = make_fvar<cpp_bin_float_50,0,Nx>(12);
  const var y = make_fvar<cpp_bin_float_50,0,0,Ny>(13);
  const var z = make_fvar<cpp_bin_float_50,0,0,0,Nz>(14);
  const var v = f(w,x,y,z);
  // Calculated from Mathematica symbolic differentiation.
  const cpp_bin_float_50 answer("1976.3196007477977177798818752904187209081211189218755");
  std::cout << std::setprecision(std::numeric_limits<cpp_bin_float_50>::digits10)
    << "mathematica   : " << answer << '\n'
    << "autodiff      : " << v.derivative(Nw,Nx,Ny,Nz) << '\n'
    << "relative error: " << std::setprecision(3) << (v.derivative(Nw,Nx,Ny,Nz)/answer-1)
    << std::endl;
  return 0;
}
/*
Output:
mathematica   : 1976.3196007477977177798818752904187209081211892188
autodiff      : 1976.3196007477977177798818752904187209081211892188
relative error: 2.67e-50
*/
```

4 Mathematics

In order for the usage of the autodiff library to make sense, a basic understanding of the mathematics will help.

4.1 Truncated Taylor Series

Basic calculus courses teach that a real analytic function $f : D \to \mathbb{R}$ is one which can be expressed as a Taylor series at a point $x_0 \in D \subseteq \mathbb{R}$:

$$f(x) = f(x_0) + f'(x_0)(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \frac{f'''(x_0)}{3!}(x - x_0)^3 + \cdots$$

One way of thinking about this form is that given the value of an analytic function $f(x_0)$ and its derivatives $f'(x_0), f''(x_0), f'''(x_0), ...$ evaluated at a point x_0, then the value of the function $f(x)$ can be obtained at any other point $x \in D$ using the above formula.

Let us make the substitution $x = x_0 + \varepsilon$ and rewrite the above equation to get:

$$f(x_0 + \varepsilon) = f(x_0) + f'(x_0)\varepsilon + \frac{f''(x_0)}{2!}\varepsilon^2 + \frac{f'''(x_0)}{3!}\varepsilon^3 + \cdots$$

Now consider ε as *an abstract algebraic entity that never acquires a numeric value*, much like one does in basic algebra with variables like x or y. For example, we can still manipulate entities like xy and $(1 + 2x + 3x^2)$ without having to assign specific numbers to them.

Using this formula, autodiff goes in the other direction. Given a general formula/algorithm for calculating $f(x_0 + \varepsilon)$, the derivatives are obtained from the coefficients of the powers of ε in the resulting computation. The general coefficient for ε^n is

$$\frac{f^{(n)}(x_0)}{n!}.$$

Thus to obtain $f^{(n)}(x_0)$, the coefficient of ε^n is multiplied by $n!$.

4.1.1 Example

Apply the above technique to calculate the derivatives of $f(x) = x^4$ at $x_0 = 2$.

The first step is to evaluate $f(x_0 + \varepsilon)$ and simply go through the calculation/algorithm, treating ε as an abstract algebraic entity:

$$\begin{aligned}
f(x_0 + \varepsilon) &= f(2 + \varepsilon) \\
&= (2 + \varepsilon)^4 \\
&= \left(4 + 4\varepsilon + \varepsilon^2\right)^2 \\
&= 16 + 32\varepsilon + 24\varepsilon^2 + 8\varepsilon^3 + \varepsilon^4.
\end{aligned}$$

Equating the powers of ε from this result with the above ε-taylor expansion yields the following equalities:

$$f(2) = 16, \qquad f'(2) = 32, \qquad \frac{f''(2)}{2!} = 24, \qquad \frac{f'''(2)}{3!} = 8, \qquad \frac{f^{(4)}(2)}{4!} = 1, \qquad \frac{f^{(5)}(2)}{5!} = 0.$$

Multiplying both sides by the respective factorials gives

$$f(2) = 16, \qquad f'(2) = 32, \qquad f''(2) = 48, \qquad f'''(2) = 48, \qquad f^{(4)}(2) = 24, \qquad f^{(5)}(2) = 0.$$

These values can be directly confirmed by the power rule applied to $f(x) = x^4$.

4.2 Arithmetic

What was essentially done above was to take a formula/algorithm for calculating $f(x_0)$ from a number x_0, and instead apply the same formula/algorithm to a polynomial $x_0 + \varepsilon$. Intermediate steps operate on values of the form

$$\mathbf{x} = x_0 + x_1\varepsilon + x_2\varepsilon^2 + \cdots + x_N\varepsilon^N$$

and the final return value is of this polynomial form as well. In other words, the normal arithmetic operators $+, -, \times, \div$ applied to numbers x are instead applied to polynomials $\mathbf{x}$. Through the overloading of C++ operators and functions, floating point data types are replaced with data types that represent these polynomials. More specifically, C++ types such as `double` are replaced with `std::array<double,N+1>`, which hold the above $N + 1$ coefficients x_i, and are wrapped in a `class` that overloads all of the arithmetic operators.

The logic of these arithmetic operators simply mirror that which is applied to polynomials. We'll look at each of the 4 arithmetic operators in detail.

4.2.1 Addition

The addition of polynomials $\mathbf{x}$ and $\mathbf{y}$ is done component-wise:

$$\begin{aligned}
\mathbf{z} = \mathbf{x} + \mathbf{y} \\
= \left(\sum_{i=0}^{N} x_i\varepsilon^i\right) + \left(\sum_{i=0}^{N} y_i\varepsilon^i\right) \\
= \sum_{i=0}^{N}(x_i + y_i)\varepsilon^i \\
z_i = x_i + y_i \qquad \text{for } i \in \{0, 1, 2, ..., N\}.
\end{aligned}$$

4.2.2 Subtraction

Subtraction follows the same form as addition:

$$\mathbf{z} = \mathbf{x} - \mathbf{y}$$

$$= \left(\sum_{i=0}^{N} x_i \varepsilon^i \right) - \left(\sum_{i=0}^{N} y_i \varepsilon^i \right)$$

$$= \sum_{i=0}^{N} (x_i - y_i) \varepsilon^i$$

$$z_i = x_i - y_i \qquad \text{for } i \in \{0, 1, 2, ..., N\}.$$

4.2.3 Multiplication

Multiplication produces higher-order terms:

$$\mathbf{z} = \mathbf{x} \times \mathbf{y}$$

$$= \left(\sum_{i=0}^{N} x_i \varepsilon^i \right) \left(\sum_{i=0}^{N} y_i \varepsilon^i \right)$$

$$= x_0 y_0 + (x_0 y_1 + x_1 y_0) \varepsilon + (x_0 y_2 + x_1 y_1 + x_2 y_0) \varepsilon^2 + \cdots + \left(\sum_{j=0}^{N} x_j y_{N-j} \right) \varepsilon^N + O\left(\varepsilon^{N+1} \right)$$

$$= \sum_{i=0}^{N} \sum_{j=0}^{i} x_j y_{i-j} \varepsilon^i + O\left(\varepsilon^{N+1} \right)$$

$$z_i = \sum_{j=0}^{i} x_j y_{i-j} \qquad \text{for } i \in \{0, 1, 2, ..., N\}.$$

In the case of multiplication, terms involving powers of ε greater than N, collectively denoted by $O\left(\varepsilon^{N+1} \right)$, are simply discarded. Fortunately, the values of z_i for $i \leq N$ do not depend on any of these discarded terms, so there is no loss of precision in the final answer. The only information that is lost are the values of higher order derivatives, which we are not interested in anyway. If we were, then we would have simply chosen a larger value of N to begin with.

4.2.4 Division

Division is not directly calculated as are the others. Instead, to find the components of $\mathbf{z} = \mathbf{x} \div \mathbf{y}$ we require that $\mathbf{x} = \mathbf{y} \times \mathbf{z}$. This yields a recursive formula for the components z_i:

$$x_i = \sum_{j=0}^{i} y_j z_{i-j}$$

$$= y_0 z_i + \sum_{j=1}^{i} y_j z_{i-j}$$

$$z_i = \frac{1}{y_0} \left(x_i - \sum_{j=1}^{i} y_j z_{i-j} \right) \qquad \text{for } i \in \{0, 1, 2, ..., N\}.$$

In the case of division, the values for z_i must be calculated sequentially, since z_i depends on the previously calculated values $z_0, z_1, ..., z_{i-1}$.

4.3 General Functions

Calling standard mathematical functions such as `log()`, `cos()`, etc. should return accurate higher order derivatives. For example, `exp(x)` may be written internally as a specific 14th-degree polynomial to approximate e^x when $0 < x < 1$. This would mean that the 15th derivative, and all higher order derivatives, would be 0, however we know that $\exp^{(15)}(x) = e^x$. How should such functions whose derivatives are known be written to provide accurate higher order derivatives? The answer again comes back to the function's Taylor series.

To simplify notation, for a given polynomial $\mathbf{x} = x_0 + x_1\varepsilon + x_2\varepsilon^2 + \cdots + x_N\varepsilon^N$ define

$$\mathbf{x}_\varepsilon = x_1\varepsilon + x_2\varepsilon^2 + \cdots + x_N\varepsilon^N = \sum_{i=1}^{N} x_i\varepsilon^i.$$

This allows for a concise expression of a general function f of $\mathbf{x}$:

$$f(\mathbf{x}) = f(x_0 + \mathbf{x}_\varepsilon)$$

$$= f(x_0) + f'(x_0)\mathbf{x}_\varepsilon + \frac{f''(x_0)}{2!}\mathbf{x}_\varepsilon^2 + \frac{f'''(x_0)}{3!}\mathbf{x}_\varepsilon^3 + \cdots + \frac{f^{(N)}(x_0)}{N!}\mathbf{x}_\varepsilon^N + O\left(\varepsilon^{N+1}\right)$$

$$= \sum_{i=0}^{N} \frac{f^{(i)}(x_0)}{i!}\mathbf{x}_\varepsilon^i + O\left(\varepsilon^{N+1}\right)$$

where ε has been substituted with $\mathbf{x}_\varepsilon$ in the ε-taylor series for $f(x)$. This form gives a recipe for calculating $f(\mathbf{x})$ in general from regular numeric calculations $f(x_0)$, $f'(x_0)$, $f''(x_0)$, ... and successive powers of the epsilon terms $\mathbf{x}_\varepsilon$.

For an application in which we are interested in up to N derivatives in x the data structure to hold this information is an $(N+1)$-element array `v` whose general element is

$$\mathtt{v[i]} = \frac{f^{(i)}(x_0)}{i!} \qquad \text{for } i \in \{0, 1, 2, ..., N\}.$$

4.4 Multiple Variables

In C++, the generalization to mixed partial derivatives with multiple independent variables is conveniently achieved with recursion. To begin to see the recursive pattern, consider a two-variable function $f(x, y)$. Since x and y are independent, they require their own independent epsilons ε_x and ε_y, respectively.

Expand $f(x, y)$ for $x = x_0 + \varepsilon_x$:

$$f(x_0 + \varepsilon_x, y) = f(x_0, y) + \frac{\partial f}{\partial x}(x_0, y)\varepsilon_x + \frac{1}{2!}\frac{\partial^2 f}{\partial x^2}(x_0, y)\varepsilon_x^2 + \frac{1}{3!}\frac{\partial^3 f}{\partial x^3}(x_0, y)\varepsilon_x^3 + \cdots + \frac{1}{M!}\frac{\partial^M f}{\partial x^M}(x_0, y)\varepsilon_x^M + O\left(\varepsilon_x^{M+1}\right)$$

$$= \sum_{i=0}^{M} \frac{1}{i!}\frac{\partial^i f}{\partial x^i}(x_0, y)\varepsilon_x^i + O\left(\varepsilon_x^{M+1}\right).$$

Next, expand $f(x_0 + \varepsilon_x, y)$ for $y = y_0 + \varepsilon_y$:

$$f(x_0 + \varepsilon_x, y_0 + \varepsilon_y) = \sum_{j=0}^{N} \frac{1}{j!}\frac{\partial^j}{\partial y^j}\left(\sum_{i=0}^{M} \varepsilon_x^i \frac{1}{i!}\frac{\partial^i f}{\partial x^i}\right)(x_0, y_0)\varepsilon_y^j + O\left(\varepsilon_x^{M+1}\right) + O\left(\varepsilon_y^{N+1}\right)$$

$$= \sum_{i=0}^{M}\sum_{j=0}^{N} \frac{1}{i!j!}\frac{\partial^{i+j} f}{\partial x^i \partial y^j}(x_0, y_0)\varepsilon_x^i\varepsilon_y^j + O\left(\varepsilon_x^{M+1}\right) + O\left(\varepsilon_y^{N+1}\right).$$

Similar to the single-variable case, for an application in which we are interested in up to M derivatives in x and N derivatives in y, the data structure to hold this information is an $(M+1) \times (N+1)$ array `v` whose element at (i, j) is

$$\mathtt{v[i][j]} = \frac{1}{i!j!}\frac{\partial^{i+j} f}{\partial x^i \partial y^j}(x_0, y_0) \qquad \text{for } (i, j) \in \{0, 1, 2, ..., M\} \times \{0, 1, 2, ..., N\}.$$

The generalization to additional independent variables follows the same pattern.

5 Migrating Code for Autodiff Compatibility

In this section, a general procedure is given for transforming an existing C++ mathematical function for compatibility with autodiff, so that single and/or multi-variable derivatives can be automatically calculated for it. The example used here is the standard Black-Scholes pricing model for pricing European options:

```cpp
#include <boost/math/constants/constants.hpp>
#include <cmath>

// Standard normal cumulative distribution function
double Phi(double x)
{
  return 0.5*std::erfc(-boost::math::constants::one_div_root_two<double>()*x);
}

enum CP { call, put };

// Black-Scholes pricing function, assuming zero annual dividend yield (q=0).
double bs_price(CP cp, double K, double S, double sigma, double tau, double r)
{
  double d1 = (std::log(S/K) + (r+sigma*sigma/2)*tau) / (sigma*std::sqrt(tau));
  double d2 = (std::log(S/K) + (r-sigma*sigma/2)*tau) / (sigma*std::sqrt(tau));
  if (cp == CP::call)
    return S*Phi(d1) - std::exp(-r*tau)*K*Phi(d2);
  else
    return std::exp(-r*tau)*K*Phi(-d2) - S*Phi(-d1);
}
```

5.1 General Rules for Transforming a C++ Function

To apply autodiff for calculating derivatives of a given function:

1. `#include <boost/math/differentiation/autodiff.hpp>`

2. `using namespace boost::math::differentiation;`

3. Select which function parameters will be variables of differentiation (aka autodiff variables). The non-variables are called constants. (Note that the use of the word *constant* here is different than the `const` C/C++ keyword. There can be `const` variables of differentiation, as well as non-`const` constants, for example, as seen below.)

4. For each function that takes at least one variable of differentiation as a parameter:

   (a) Replace the data type (e.g. `double`) of each variable parameter with a generic template type. For best performance, make it a reference. For best practice, also make it `const` if appropriate. Declare a separate template type for each variable parameter.

   (b) Replace the return value with `promote<T1,T2,T3>` where T1, T2, T3 are the template types created in the previous step. There can be any number of types. (If there is only 1 type T then the return type can simply be T instead of `promote<T>`).

   (c) Update the data type of all temporary variables that depend on an autodiff variable by `auto` (easiest) or where that is not possible, then use `promote<...>` to combine the data types of the autodiff variables in the intermediate calculation. Omitting any necessary data types from `promote<>` can lead to incorrect results, which can be avoided by using `auto`.

   (d) Remove the namespace prefix from mathematical function calls applied to autodiff variables, to allow for overloaded functions to be called via argument-dependent lookup (ADL). This requires that the function exist in the autodiff library of overloaded functions (E.g. `exp`, `log`, `erfc`, etc.). If it doesn't, then an appropriate alternative must be made available.

5. Each variable of differentiation must be separately declared and initialized by calling the `make_fvar<...>()` template function. Say there are M such variables. For each variable:

   (a) Select the underlying `root_type` for the variable. E.g. `float`, `double`, `long double`, `boost::multiprecision::cpp_bin_float_50`, etc. This is always the first template parameter.

(b) Select what the highest order derivative N_i will be calculated for it, where $i \in \{1, 2, ..., M\}$. That, and all lower order derivatives will be available in the result. This must be a compile-time, or `constexpr` value. Specify this value as the i^{th} template parameter to `make_fvar`, where the 0^{th} parameter is the `root_type`, and all parameters in between set to 0. **The number of template parameters to `make_fvar` is how autodiff distinguishes independent variables from one another.** If two independent variables are declared with the same number of template parameters, and they are used in the same calculation, then the results may be non-sensical (unless you deliberately understand the mathematical consequences of this.)

(c) Set the actual value of the parameter as the function parameter. This is stored internally as `root_type` and may be determined at run-time.

For example, to set the 3^{rd} variable to 2.0, and calculate up to $N_3 = 5$ derivatives for it, the initialization would be `auto x3 = make_fvar<double,0,0,5>(2.0);`

6. Call the function, and store the result as type `auto` (easiest) or of type `promote<T1,T2,T3,...>`:
   `auto y = f(x1, x2, x3, ...);`

7. Call `y.derivative(...)` to obtain all mixed partial derivatives, with the following requirements:

   (a) `derivative(...)` must be called with exactly M parameters (one for each variable.)

   (b) The i^{th} parameter to `derivative(...)` (where $i = 1$ corresponds to the first parameter) must be of type `size_t` with a maximum value of N_i.

8. For best performance, compile with the C++17 flag, or highest available, with C++11 at a minimum.

5.2 Example: Black-Scholes Option Pricing with Automatic Greeks

The above procedure is applied below to the Black-Scholes pricing function `bs_price()`. Of the 6 parameters to `bs_price()`, the first 2 (call/put `cp` and strike price K) are constant, since no derivatives are calculated with respect to them, and the remaining 4 are autodiff variables. A `main()` function is added to demonstrate how to initialize the $M = 4$ independent variables with `make_fvar<>()`, invoke the function `bs_price()`, and retrieve its derivatives:

```cpp
#include <boost/math/constants/constants.hpp>
#include <cmath>
#include <iostream>

#include <boost/math/differentiation/autodiff.hpp>
using namespace boost::math::differentiation;

// Standard normal cumulative distribution function
template<typename T>
T Phi(const T& x)
{
  using std::erfc;
  return 0.5*erfc(-boost::math::constants::one_div_root_two<T>()*x);
}

enum CP { call, put };

// Black-Scholes pricing function, assuming zero annual dividend yield (q=0).
template<typename T1, typename T2, typename T3, typename T4>
promote<T1,T2,T3,T4>
  bs_price(CP cp, double K, const T1& S, const T2& sigma, const T3& tau, const T4& r)
{
  using namespace std;
  auto d1 = (log(S/K) + (r+sigma*sigma/2)*tau) / (sigma*sqrt(tau));
  auto d2 = (log(S/K) + (r-sigma*sigma/2)*tau) / (sigma*sqrt(tau));
  if (cp == CP::call)
    return S*Phi(d1) - exp(-r*tau)*K*Phi(d2);
  else
    return exp(-r*tau)*K*Phi(-d2) - S*Phi(-d1);
}

int main()
```

```
{
  double K     = 100.0; // Strike price
  auto S       = make_fvar<double,3>(105); // Spot price
  auto sigma   = make_fvar<double,0,3>(5); // Volatility
  auto tau     = make_fvar<double,0,0,1>(30.0/365); // 30 days to expiration
  auto r       = make_fvar<double,0,0,0,1>(1.25/100); // 1.25% interest rate
  auto call_price = bs_price(CP::call, K, S, sigma, tau, r);
  auto put_price  = bs_price(CP::put , K, S, sigma, tau, r);

  std::cout << std::setprecision(std::numeric_limits<double>::digits10)
    << "call price   = " <<  call_price.derivative(0,0,0,0) << '\n'
    << "put  price   = " <<   put_price.derivative(0,0,0,0) << '\n'
    << "\n## First-order Greeks\n"
    << "call delta   = " <<  call_price.derivative(1,0,0,0) << '\n'
    << "call vega    = " <<  call_price.derivative(0,1,0,0) << '\n'
    << "call theta   = " << -call_price.derivative(0,0,1,0) << '\n' // minus sign due to
    << "call rho     = " <<  call_price.derivative(0,0,0,1) << '\n' //     tau = T-time
    << '\n'
    << "put delta    = " <<   put_price.derivative(1,0,0,0) << '\n'
    << "put vega     = " <<   put_price.derivative(0,1,0,0) << '\n'
    << "put theta    = " <<  -put_price.derivative(0,0,1,0) << '\n'
    << "put rho      = " <<   put_price.derivative(0,0,0,1) << '\n'
    << "\n## Second-order Greeks (same for call and put)\n"
    << "call gamma   = " <<  call_price.derivative(2,0,0,0) << '\n'
    << "put  gamma   = " <<   put_price.derivative(2,0,0,0) << '\n'
    << "call vanna   = " <<  call_price.derivative(1,1,0,0) << '\n'
    << "put  vanna   = " <<   put_price.derivative(1,1,0,0) << '\n'
    << "call charm   = " << -call_price.derivative(1,0,1,0) << '\n'
    << "put  charm   = " <<  -put_price.derivative(1,0,1,0) << '\n'
    << "call vomma   = " <<  call_price.derivative(0,2,0,0) << '\n'
    << "put  vomma   = " <<   put_price.derivative(0,2,0,0) << '\n'
    << "call veta    = " <<  call_price.derivative(0,1,1,0) << '\n'
    << "put  veta    = " <<   put_price.derivative(0,1,1,0) << '\n'
    << "\n## Third-order Greeks (same for call and put)\n"
    << "call speed   = " <<  call_price.derivative(3,0,0,0) << '\n'
    << "put  speed   = " <<   put_price.derivative(3,0,0,0) << '\n'
    << "call zomma   = " <<  call_price.derivative(2,1,0,0) << '\n'
    << "put  zomma   = " <<   put_price.derivative(2,1,0,0) << '\n'
    << "call color   = " <<  call_price.derivative(2,0,1,0) << '\n'
    << "put  color   = " <<   put_price.derivative(2,0,1,0) << '\n'
    << "call ultima  = " <<  call_price.derivative(0,3,0,0) << '\n'
    << "put  ultima  = " <<   put_price.derivative(0,3,0,0) << '\n';
  return 0;
}
```

Note the correspondence between the number of template parameters to `make_fvar<...>` and the order of parameters to `derivative(...)`. The 4 independent variables are initialized as:

```
auto S       = make_fvar<double,3>(105); // Spot price
auto sigma   = make_fvar<double,0,3>(5); // Volatility
auto tau     = make_fvar<double,0,0,1>(30.0/365); // 30 days to expiration
auto r       = make_fvar<double,0,0,0,1>(1.25/100); // 1.25% interest rate
```

This means the variable `S` corresponds to variable 1, `sigma` to 2, `tau` to 3, and `r` to 4. The maximum number of derivatives to be calculated for each are $N_1 = 3$, $N_2 = 3$, $N_3 = 1$, and $N_4 = 1$, respectively.

Taking for example the calculation of color:

```
<< "call color   = " <<  call_price.derivative(2,0,1,0) << '\n'
```

This returns the 2nd partial derivative of `call_price` with respect to variable 1 (`S`) and the 1st partial derivative with respect to variable 3 (`tau`).

5.3 Advantages of Automatic Differentiation

This example also provides an opportunity to point out some of the advantages of using autodiff:

- Elimination of code redundancy. The existence of N separate functions to calculate derivatives is a form of code redundancy, with all the liabilities that come with it:

  - Changes to one function require N additional changes to other functions. Consider how much larger and inter-dependent the above code base would be if a separate function were written for each Greek value.

  - Dependencies upon a derivative function for a different purpose will break when changes are made to the original function. These dependencies won't exist if such derivative functions don't exist.

  - Code bloat, reducing conceptual integrity. Control over the evolution of code is easier/safer when the code base is smaller and able to be intuitively grasped.

- Accuracy of derivatives over finite difference methods. Finite difference methods always include a Δx free variable that must be carefully chosen for each application. If Δx is too small, then numerical errors become large. If Δx is too large, then mathematical errors become large. With autodiff, there are no free variables to set and the accuracy of the answer is generally superior to finite difference methods even with the best choice of Δx.

6 Function Writing Guidelines

At a high level there is one fairly simple principle, loosely and intuitively speaking, to writing functions for which autodiff can effectively calculate derivatives:

Autodiff Function Principle (AFP)

A function whose branches in logic correspond to piecewise analytic calculations over non-singleton intervals, with smooth transitions between the intervals, and is free of indeterminate forms in the calculated value and higher order derivatives, will work fine with autodiff.

Stating this with greater mathematical rigor can be done. However what seems to be more practical, in this case, is to give examples and categories of examples of what works, what doesn't, and how to remedy some of the common problems that may be encountered. That is the approach taken here.

6.1 Example 1: $f(x) = \max(0, x)$

One potential implementation of $f(x) = \max(0, x)$ is:

```
template<typename T>
T f(const T& x)
{
    return 0 < x ? x : 0;
}
```

Though this is consistent with Section 5, there are two problems with it:

1. `f(nan) = 0`. This problem is independent of autodiff, but is worth addressing anyway. If there is an indeterminate form that arises within a calculation and is input into f, then it gets "covered up" by this implementation leading to an unknowingly incorrect result. Better for functions in general to propagate NaN values, so that the user knows something went wrong and doesn't rely on an incorrect result, and likewise the developer can track down where the NaN originated from and remedy it.

2. $f'(0) = 0$ when autodiff is applied. This is because f returns 0 as a constant when x==0, wiping out any of the derivatives (or sensitivities) that x was holding as an autodiff variable. Instead, let us apply the AFP and identify the two intervals over which f is defined: $(-\infty, 0] \cup (0, \infty)$. Though the function itself is not analytic at $x = 0$, we can attempt somewhat to smooth out this transition point by averaging the calculation of $f(x)$ at $x = 0$ from each interval. If $x < 0$ then the result is simply 0, and if $0 < x$ then the result is x. The average is $\frac{1}{2}(0 + x)$ which will allow autodiff to calculate $f'(0) = \frac{1}{2}$. This is a more reasonable answer.

A better implementation that resolves both issues is:

```
template<typename T>
T f(const T& x)
{
    if (x < 0)
        return 0;
    else if (x == 0)
        return 0.5*x;
    else
        return x;
}
```

6.2   Example 2: $f(x) = \mathrm{sinc}(x)$

The definition of $\mathrm{sinc} : \mathbb{R} \to \mathbb{R}$ is

$$\mathrm{sinc}(x) = \begin{cases} 1 & \text{if } x = 0 \\ \frac{\sin(x)}{x} & \text{otherwise.} \end{cases}$$

A potential implementation is:

```
template<typename T>
T sinc(const T& x)
{
    using std::sin;
    return x == 0 ? 1 : sin(x) / x;
}
```

Though this is again consistent with Section 5, and returns correct non-derivative values, it returns a constant when `x==0` thereby losing all derivative information contained in `x` and contributions from sinc. For example, $\mathrm{sinc}''(0) = -\frac{1}{3}$, however `y.derivative(2) == 0` when `y = sinc(make_fvar<double,2>(0))` using the above incorrect implementation. Applying the AFP, the intervals upon which separate branches of logic are applied are $(-\infty, 0) \cup [0, 0] \cup (0, \infty)$. The violation occurs due to the singleton interval $[0, 0]$, even though a constant function of 1 is technically analytic. The remedy is to define a custom sinc overload and add it to the autodiff library. This has been done. Mathematically, it is well-defined and free of indeterminate forms, as is the 3^{rd} expression in the equalities

$$\frac{1}{x}\sin(x) = \frac{1}{x}\sum_{n=0}^{\infty}\frac{(-1)^n}{(2n+1)!}x^{2n+1} = \sum_{n=0}^{\infty}\frac{(-1)^n}{(2n+1)!}x^{2n}.$$

The autodiff library contains helper functions to help write function overloads when the derivatives of a function are known. This is an advanced feature and documentation for this may be added at a later time.

For now, it is worth understanding the ways in which indeterminate forms can occur within a mathematical calculation, and avoid them when possible by rewriting the function. Table 1 compares 3 types of indeterminate forms. Assume the product `a*b` is a positive finite value.

	$f(x) = \left(\frac{a}{x}\right) \times (bx^2)$	$g(x) = \left(\frac{a}{x}\right) \times (bx)$	$h(x) = \left(\frac{a}{x^2}\right) \times (bx)$
Mathematical Limit	$\lim_{x \to 0} f(x) = 0$	$\lim_{x \to 0} g(x) = ab$	$\lim_{x \to 0} h(x) = \infty$
Floating Point Arithmetic	`f(0) = inf*0 = nan`	`g(0) = inf*0 = nan`	`h(0) = inf*0 = nan`

Table 1: Automatic differentiation does not compute limits. Indeterminate forms must be simplified manually. (These cases are not meant to be exhaustive.)

Indeterminate forms result in NaN values within a calculation. Mathematically, if they occur at locally isolated points, then we generally prefer the mathematical limit as the result, even if it is infinite. As demonstrated in Table 1, depending upon the nature of the indeterminate form, the mathematical limit can be 0 (no matter the values of a or b), or ab, or ∞, but these 3 cases cannot be distinguished by the floating point result of nan. Floating point arithmetic does not perform limits (directly), and neither does the autodiff library. Thus it is up to the diligence of the developer to keep a watchful eye over where indeterminate forms can arise.

6.3 Example 3: $f(x) = \sqrt{x}$ and $f'(0) = \infty$

When working with functions that have infinite higher order derivatives, this can very quickly result in nans in higher order derivatives as the computation progresses, as `inf-inf`, `inf/inf`, and `0*inf` result in `nan`. See Table 2 for an example.

$f(x)$	$f(0)$	$f'(0)$	$f''(0)$	$f'''(0)$
sqrt(x)	0	inf	-inf	inf
sqr(sqrt(x)+1)	1	inf	nan	nan
x+2*sqrt(x)+1	1	inf	-inf	inf

Table 2: Indeterminate forms in higher order derivatives. `sqr(x) == x*x`.

Calling the autodiff-overloaded implementation of $f(x) = \sqrt{x}$ at the value `x==0` results in the 1$^{\text{st}}$ row (after the header row) of Table 2, as is mathematically correct. The 2$^{\text{nd}}$ row shows $f(x) = (\sqrt{x} + 1)^2$ resulting in `nan` values for $f''(0)$ and all higher order derivatives. This is due to the internal arithmetic in which `inf` is added to `-inf` during the squaring, resulting in a `nan` value for $f''(0)$ and all higher orders. This is typical of `inf` values in autodiff. Where they show up, they are correct, however they can quickly introduce `nan` values into the computation upon the addition of oppositely signed `inf` values, division by `inf`, or multiplication by 0. It is worth noting that the infection of `nan` only spreads upward in the order of derivatives, since lower orders do not depend upon higher orders (which is also why dropping higher order terms in an autodiff computation does not result in any loss of precision for lower order terms.)

The resolution in this case is to manually perform the squaring in the computation, replacing the 2$^{\text{nd}}$ row with the 3$^{\text{rd}}$: $f(x) = x + 2\sqrt{x} + 1$. Though mathematically equivalent, it allows autodiff to avoid `nan` values since $\sqrt{x}$ is more "isolated" in the computation. That is, the `inf` values that unavoidably show up in the derivatives of `sqrt(x)` for `x==0` do not have the chance to interact with other `inf` values as with the squaring.

6.4 Summary

The AFP gives a high-level unified guiding principle for writing C++ template functions that autodiff can effectively evaluate derivatives for.

Examples have been given to illustrate some common items to avoid doing:

1. It is not enough for functions to be piecewise continuous. On boundary points between intervals, return the average expression of both intervals, rather than just one of them. Example: $\max(0, x)$ at $x = 0$.

2. Avoid returning individual constant values (e.g. $\text{sinc}(0) = 1$.) Values must be computed uniformly along with other values in its local interval. If that is not possible, then the function must be overloaded to compute the derivatives precisely.

3. Avoid intermediate indeterminate values in both the value ($\text{sinc}(x)$ at $x = 0$) and derivatives ($(\sqrt{x} + 1)^2$ at $x = 0$). Try to isolate expressions that may contain infinite values/derivatives so that they do not introduce NaN values into the computation.

This is not an exhaustive list, and will be added to as examples are encountered.

7 Acknowledgments

References

[1] `https://en.wikipedia.org/wiki/Automatic_differentiation`

[2] Andreas Griewank, Andrea Walther. *Evaluating Derivatives*. SIAM, 2nd ed. 2008.